

Intentional AI Leadership + Scaling

Justworks.org Capacity Building Workshop — Pre-Read

Before we meet. AI is becoming part of how organizations communicate, analyze information, manage relationships, document knowledge, and get work done. But this workshop isn't about becoming an AI expert or chasing the newest tool. It's about developing a way of thinking that helps us answer a more useful question:

How can we create more capacity for the work that matters?

We'll explore how AI can help us **simplify, augment, automate, and scale**—while keeping people, judgment, and organizational context at the center. This pre-read is the primer. The external links throughout are optional deep dives; you don't need to read them all before we meet.

[Start with the problem, not the tool](#)

[Same concepts, different buzzwords: a cross-tool decoder](#)

[From "chatbot" to organizational capability](#)

[Data, information, and institutional knowledge are different](#)

[AI as a search box vs. an organizational knowledge layer](#)

[What makes organizational knowledge trustworthy?](#)

[Additional Links](#)

[AI and social impact](#)

[Responsible AI policies](#)

[Google for Nonprofits](#)

[Buzzwords](#)

Start with the problem, not the tool

It's easy to start with "Can AI do this?" We want to pause and ask a better first question: **What are we actually trying to accomplish?** Then work through it in order:

- 1. Understand the problem.** What are we trying to accomplish?
- 2. Understand the current workflow.** How does work happen today?
- 3. Find friction.** Where are we losing time, creating errors, duplicating work, or bottlenecks?
- 4. Find the root cause.** Why does that friction exist?
- 5. Simplify.** Can we eliminate or redesign part of the process?
- 6. Consider AI.** Could AI augment or automate what's left?
- 7. Humans in the loop.** Where are judgment, relationships, context, and accountability essential?

A useful rule: the best AI solution might not be an AI solution. Sometimes the answer is a better process, clearer documentation, better data, or simply stopping work that doesn't need to happen.

Same concepts, different buzzwords: a cross-tool decoder

The best tool isn't necessarily the model with the best demo. It's the tool that fits the problem, workflow, data, people, and risk.

Every AI tool markets its own features with its own vocabulary—Gems, Custom GPTs, connectors, skills, artifacts, tasks, memory. It can feel like each tool is a completely different world. Underneath, they're not. Every one of these tools is built on the same core idea: **machine learning systems that recognize patterns in data and predict the most useful next output**—the next word, the next classification, the next suggested action. The branded features are just different wrappers around a small set of repeatable capabilities. Learn the pattern, and almost any tool becomes learnable.

Concept	What it actually does	What different tools call it
The engine	Recognizes patterns in language and data and predicts the most useful next output. Learns from structured & unstructured data and reinforcement	Machine learning / large language models — sitting under every tool
Context	The information you give it right now: your prompt, attached files, and the current conversation	Prompt, context window, attachments (all tools)
Memory*	What it retains about you across conversations so you don't have to repeat yourself	ChatGPT Memory; Gemini personal context; Claude memory
Reusable custom assistant	A saved assistant preloaded with instructions, a role, and reference knowledge	Gemini Gems ; ChatGPT Custom GPTs ; Claude Projects
Connections to your systems	Links the AI to your other tools and data (Drive, email, CRM) so it can read or act	ChatGPT connectors / Apps; Gemini connected apps & extensions; Claude connectors — often built on MCP
Repeatable tasks	Packaged, on-demand instructions the AI pulls in for a specific kind of task	Claude Skills (and similar capability packs)
A workspace for outputs	A canvas or document where generated content lives and can be edited, kept, and revisited	Claude Artifacts; ChatGPT Canvas; Gemini Canvas
Automations	The AI runs on a schedule or trigger instead of only when you ask	ChatGPT Tasks; Claude scheduled tasks

*Memory works differently on each platform

All three platforms remember you, with different mechanics. ChatGPT's Memory is largely

automatic; the model decides what to save from your conversations. Claude's approach is more explicit and document-driven, since Project files get read every time and you control exactly what persists. Gemini's memory is really three things stacked together: Saved Info (facts you tell it), references to past chats, and context pulled from connected Google apps like Gmail and Drive. That last part is the real differentiator. The catch for teams: none of this memory transfers between vendors, so whichever platform accumulates your context becomes harder to leave.

[ChatGPT vs Claude vs Gemini: Feature & Terminology Map \(2026\)](#)

[FAQ by simpliaxis : Claude vs ChatGPT vs Gemini: Features, Pricing, Use Cases & Future Outlook](#)

Notice the pattern. Most "new" features are one of a handful of ideas: give it context, let it remember, save a reusable version, connect it to your systems, package a repeatable skill, give outputs a home, or let it run automatically. So when a tool announces something new, the useful question isn't "how do I learn this from scratch?"—it's "which of these patterns is this?" That's the same problem-first instinct we'll use all workshop: **see the underlying pattern, not the shiny label.**

The Terminology Translation Table [\(link\)](#)

What it does	ChatGPT (OpenAI)	Claude (Anthropic)	Gemini (Google)
Persistent workspace: a folder of chats with shared files and instructions for one client, campaign, or initiative	Projects	Projects	Gems (closest thing, see note below)
Custom reusable assistant: a saved persona with its own instructions and knowledge files, like a “Brand Voice Editor”	Custom GPTs	Projects with custom instructions, plus Skills	Gems
Reusable instruction modules: packaged how-to instructions the AI applies to a recurring task	Skills / GPT instructions	Skills	Gem instructions
Side-by-side editing pane for drafting documents, code, or prototypes	Canvas	Artifacts	Canvas
Long-running research agent that produces a cited report	Deep Research	Research	Deep Research
Cross-chat memory (remembers facts about you between conversations)	Memory (saved memories)	Memory	Saved Info + Personal Context (past chats)
Standing preferences about tone, format, and context	Custom Instructions	Preferences + Styles	Saved Info
Scheduled automations (daily briefings, weekly reports)	Tasks	Scheduled Tasks (via Cowork)	Scheduled Actions
Third-party integrations with your CRM, drive, calendar, etc.	Connectors + Apps (and GPT Actions)	Connectors (built on MCP)	Apps / Workspace integration (Gmail, Docs, Drive)
Agentic browsing, where the AI operates a browser or desktop for you	Agent Mode (formerly Operator)	Claude for Chrome + Cowork	Agent Mode
Coding agent that works in your codebase or terminal	Codex	Claude Code	Gemini CLI + Jules
Marketplace / distributing custom assistants to other people	GPT Store	Org-level sharing on Team/Enterprise; no public store	Gem link sharing; no public store
Voice conversation	Voice Mode (Advanced Voice)	Voice mode (mobile)	Gemini Live
Underlying model family	GPT series (GPT-5.x)	Claude (Opus, Sonnet, Haiku tiers)	Gemini (Pro and Flash tiers)

THE EXECUTIVE'S AI TEAM (NOV 2025)



GEMINI 3 (The Deep Researcher)
- Handles the "Junk Drawer" of messy data (audio, video, PDFs).
Your Lead Scientist.



CHATGPT 5.1 (The Generalist)
- Master of comms, quick tasks, and structured data.
Your Chief of Staff.



CLAUDE OPUS 4.5 (The Thinker)
- For complex coding, logic, and system architecture.
Your Senior Engineer.

DIFFERENT TOOLS FOR DIFFERENT JOBS. BUILD YOUR ROSTER.

^ [Your AI Tools Aren't Competing. They're Applying for Different Jobs](#)

AI STRENGTHS COMPARISON: CHATGPT vs CLAUDE vs GEMINI

EXPLORING KEY ADVANTAGES AND CAPABILITIES

STRENGTHS CATEGORY	CHATGPT (OPENAI)	CLAUDE (ANTHROPIC)	GEMINI (GOOGLE)
OVERALL VERSATILITY & REASONING	 Broad versatility, strong reasoning (GPT-4), reasoning over varied tasks.	 Excellent creative and analytical writing, strong understanding of nuance.	 Deep logical reasoning, strong general knowledge, multi-step problem solving.
CONTENT GENERATION (Writing, Coding)	 Powerful content generation, creative writing, excellent coding assistant.	 Nuanced prose, high-quality professional writing, less prone to errors.	 Competitive code generation, powerful multimodal content, diverse writing styles.
MULTIMODAL CAPABILITIES	 Integrated text and image creation (DALL-E 3), web browsing, data analysis.	 Exceptional text understanding, summarization, long document analysis. Vision inputs (C3).	 NATIVE MULTIMODAL from scratch (Text, Image, Video, Audio understanding).
CONTEXT WINDOW & COMPREHENSION	 Standard context window, good instruction following.	 MASSIVE CONTEXT WINDOW (200k+ tokens), handles full books & large data.	 HUGE CONTEXT WINDOW (Up to 1M+ tokens in 1.5 Pro). Complex data analysis.
INTEGRATION & ECOSYSTEM	 OpenAI platform, custom GPTs, wide developer API usage.	 Primarily via API for businesses, clean web interface.	 DEEP GOOGLE INTEGRATION (Workspace: Docs, Gmail; Search; Android)

^ [ChatGPT vs Claude vs Gemini: I Used All Three for 60 Days — Here's the Honest Verdict](#)

CLAUDE vs CHATGPT vs GEMINI

WHAT EACH IS · WHEN TO USE · WHERE IT FITS BEST



CLAUDE

WHAT IT IS

- Anthropic's AI assistant.
- Built for depth and accuracy.
- Pro \$20/mo · Max from \$100/mo.

WHEN TO USE

- High-stakes, accuracy-critical work.
- Long documents (up to 1M tokens).
- Coding, contracts, financial analysis.

USE CASES

- Contract review and analysis.
- Financial models and reports.
- Code review and debugging.

STRENGTHS

- Handles long documents and complex analysis well.
- Less likely to make things up when unsure.
- Projects load business context.

WEAKNESSES

- No native image generation.
- Tone runs measured and deliberate.
- Tighter usage limits at lower tiers.

WHERE TO START

Set up a Project with your business context, and Claude arrives briefed every session.



CHATGPT

WHAT IT IS

- OpenAI's AI assistant.
- The most widely used AI tool.
- Plus \$20/mo · Pro from \$100/mo.

WHEN TO USE

- Speed and ideation work.
- When you need image generation.
- Building custom GPTs for recurring tasks.

USE CASES

- Drafting posts, emails, marketing copy.
- Summarizing PDFs and articles.
- Custom GPTs for weekly reports.

STRENGTHS

- Fast for everyday tasks and quick ideas.
- Largest plugin ecosystem (GPT Store).
- DALL-E image generation built in.

WEAKNESSES

- Confidently wrong sometimes.
- Higher hallucination rate than Claude.
- Tends to agree with you instead of pushing back.

WHERE TO START

Build a custom GPT for your weekly business reporting and run it every Monday.



GEMINI

WHAT IT IS

- Google's AI assistant.
- Lives inside Google Workspace.
- Pro \$19.99/mo · Ultra from \$99.99/mo.

WHEN TO USE

- When your business runs on Google Workspace.
- Real-time Google Search.
- Multimodal work (video, audio, images).

USE CASES

- Summarize emails and Docs.
- Build presentations in Slides.
- Deep Research across Gmail, Drive, Meet.

STRENGTHS

- Deepest Workspace integration.
- Native image AND video generation (Imagen + Veo).
- Deep Research across web, Gmail, Drive, and Meet.

WEAKNESSES

- Best inside Google ecosystem.
- Stability issues — chats can hang or respond inconsistently.
- New compute-based limits can cap heavy users.

WHERE TO START

Set up a Gem with your business context so Gemini brings it into every Workspace task.



FOLLOW **LEYSAN GILFANOVA** | PRACTICAL AI FOR BUSINESS OWNERS

^ [Linkedin Post](#)

From "chatbot" to organizational capability

One of the biggest shifts happening with AI is the move from "I ask AI a question" to "AI can work with the information and systems I already use." Consider four patterns:

- **Level 1 — Bring information to AI.** Upload a spreadsheet, PDF, document, transcript, or notes, then ask AI to summarize, analyze, or transform it. *Good for exploration and one-off work.*
- **Level 2 — Connect AI to organizational information.** AI accesses information from Google Drive, Microsoft 365, Salesforce, HubSpot, or other approved applications. *Good for recurring knowledge retrieval and workflows.*
- **Level 3 — Connect systems.** APIs, integrations, databases, data pipelines, and tools like MCP let systems exchange information or take actions. *Good for repeatable workflows and automation.*
- **Level 4 — Build a reliable data foundation.** Information flows through source systems → data pipelines → governed data → analytics / AI. *Good for scale, reporting, automation, and more sophisticated AI applications.*

The important lesson: these aren't levels every organization needs to "graduate" through. More technical does not automatically mean better. The right question is always: what's the simplest, safest, and most reliable way to solve the problem?

Data, information, and institutional knowledge are different

When we think about impact, the story rarely lives in one database. Evidence takes many forms.

What we're trying to understand	Examples of evidence
Who	Participants, donors, founders, partners
What we invested	Money, grants, resources
How much effort	Staff time, coaching hours, volunteer hours
What Happened	Events, workshops, coaching, referrals
What was said	Emails, interviews, surveys, transcripts
What we know	Operating procedure, notes, lessons learned, institutional knowledge
What we produced	Services delivered, people served, applications completed
What changed	Revenue, employment, confidence, access to capital, behavior
Why it matters	Context, interpretation, community knowledge, human judgment

This evidence lives across CRM, spreadsheets, accounting, email, documents, surveys, transcripts, and people's expertise. **No single system necessarily contains the whole story.** AI can help connect and synthesize different sources of evidence. But AI does not automatically know which source is authoritative, whether information is current, whether a metric is meaningful, whether a correlation represents impact, what context is missing, or what decision should be made. That's where human judgment comes in.

Example differences:

Data: Participant #382 attended three workshops.

Information: Participant #382 attended 3 of 5 available workshops.

Knowledge: Participants who attend at least three workshops are considered "engaged."

Institutional knowledge: "We consider someone engaged after their third workshop because that's historically when we see participants begin using the coaching resources. Orientation doesn't count."

That last example carries context, reasoning, organizational practice, history, and judgment—and it usually lives in **people**, not systems.

AI as a search box vs. an organizational knowledge layer

Imagine an employee asks: "How do I run our monthly participant outreach?"

AI as a search box: "Here are six documents that might be relevant." The employee still has to figure out what's correct.

AI as an organizational knowledge layer: "Here is the current process based on the organization's approved documentation. The current operating procedure says X, while a more recent program update changed step 3 to Y. I've flagged the discrepancy because the sources conflict. Sources: standard operating procedures, Program Update, CRM Field Guide."

The second is a far more powerful vision—but it requires something first: **good organizational knowledge**. AI cannot reliably create organizational truth from a collection of contradictory or outdated information.

Why the organizational knowledge layer matters

The difference between those two answers isn't the model—it's what the model is grounded in. When AI draws on your trusted, maintained sources instead of whatever someone happened to paste in, answers become consistent (everyone gets the same guidance), checkable (it can cite where the answer came from), and they get better over time as the underlying knowledge is maintained. That's the shift from a clever individual tool to a shared organizational capability.

Here's the catch, and it's the whole reason this matters for us: AI doesn't create organizational truth—it exposes the gaps you already live with. Contradictory documents, the step that changed but never got updated, the exception nobody wrote down, the process that lives only in one person's head. A person reading your documentation fills those gaps with judgment and a quick Slack message to a colleague. AI can't. It will confidently answer from whatever it's given—so "human-ready" documentation is not automatically "AI-ready," and garbage in still means garbage out. For nonprofits, where so much institutional knowledge walks out the door

with staff turnover, building this layer is also how you stop losing what you know.

- [How Effective Knowledge Management Can Enhance Nonprofit Impact](#) — KM Institute. People, processes, tools, and strategy for capturing knowledge before turnover erases it.
- [Garbage In, Garbage Out: How to Stop Your AI from Hallucinating](#) — Shelf. Plain explanation of why data quality drives whether you can trust AI's answers.

What makes organizational knowledge trustworthy?

Before connecting AI to organizational knowledge, ask:

Source quality — Is this information authoritative?

Accuracy — Is it correct?

Freshness — Is it still current?

Completeness — Are exceptions or edge cases missing?

Consistency — Do different sources agree?

Permissions — Should this person have access?

Privacy — Does it contain sensitive or confidential information?

Citations — Can we trace an answer back to its source?

Ownership — Who is responsible for the knowledge?

Maintenance — What happens when the process changes?

Additional Links

AI and social impact

- [Four Questions to Frame AI Strategy in Purpose and Values — Stanford Social Innovation Review](#). Especially relevant to this workshop: staff are experimenting, boards and funders are asking about strategy, and organizations aren't always sure how to move from scattered experimentation toward intentional adoption.

- [Leading Social Impact Strategy: Harnessing the Power of Purpose and AI — Stanford HAI](#)
Takeaway: AI strategy should be grounded in purpose, organizational capacity, and values—not technology alone.

- [NTEN — AI for Nonprofits Resource Hub](#). A nonprofit-specific collection covering AI governance, data and IT governance, privacy, tool evaluation, AI policies, responsible implementation, and human-centered, equitable AI. A resource to **bookmark** rather than read cover-to-cover.

Responsible AI policies

- [Getting Started on a Responsible AI Use Policy for Nonprofits — Candid](#). Emphasizes agreements that let staff experiment while protecting people, data, trust, and accountability.
- [What if my organization has no AI use policy? — Candid](#)
- [Responsible AI Use Policies: A Guide for Nonprofits — GlobalGiving](#)

Google for Nonprofits

- [Responsible AI for nonprofits](#) — a practical overview of accuracy, bias, privacy, human review, transparency, and appropriate use.
- [Nonprofit Gemini Prompt Library](#) — examples across fundraising, communications, qualitative research, and impact storytelling.

Buzzword Dictionary

- **Artificial Intelligence.** A broad field of computing focused on systems that perform tasks associated with human intelligence—recognizing patterns, making predictions, understanding language, or taking actions.
- **Machine learning.** The approach behind most modern AI: instead of being programmed with explicit rules, a system learns patterns from large amounts of data, then uses those patterns to make predictions or recognize new examples.
- **Generative AI.** AI that generates or transforms content: text, images, audio, video, code, summaries, or structured information. ChatGPT, Claude, and Gemini are generative AI applications.
- **Large Language Model (LLM).** A type of AI model trained to recognize and generate patterns in language. LLMs are the underlying technology behind many generative AI tools.
- **AI assistant / copilot.** AI is embedded into a person's workflow to help them accomplish a task—for example, "Draft a donor email based on this donor's history" or "Summarize the key decisions from this meeting."
- **AI-enabled workflow.** A process intentionally designed so that people and AI work together. For example: CRM data → AI identifies potential donor segments → AI drafts personalized outreach → human reviews → campaign is sent → results feed back into the process.
- **Agentic AI:** The type of AI that involves AI agents. AI agents are systems that combine advanced AI models with access to your everyday tools like email, calendars, or documents. This allows them to take action on your behalf and under your control. For nonprofits, this might look like: an AI agent that monitors incoming volunteer inquiries, checks the organization's central calendar for orientation dates, and automatically sends invitation links to prospective volunteers while updating a central contact list.
- **"AI native"** refers to a product, service, or company built from the ground up with artificial intelligence as its core foundation, rather than having AI bolted on later as an auxiliary feature. If the AI component were removed from an AI-native system, the

system would cease to be useful.

AI doesn't have to know everything. It can be connected to the information and tools it needs.

- **API.** A way for one system to request information or functionality from another. "Salesforce, give me these participant records."
- **Database.** A structured system for storing and retrieving information. "Find all participants who attended more than three workshops." SQL is a common language for querying databases.
- **Data pipeline.** A repeatable process that moves and transforms information between systems. Salesforce → pipeline → data warehouse → reporting.
- **Connector.** A prebuilt connection between an application and another tool. AI tool ↔ G Drive.
- **MCP.** Model Context Protocol is a standard for connecting AI applications with external tools, systems, and sources of information. It's the plumbing behind many of the "connectors" in the decoder above.

